



MOFFETT AI

Introduction to Sparse Computing Whitepaper

MOFFETT AI
AUGUST 2022

Copyright Notice

The copyright of this white paper belongs to Moffett AI and It is protected by law. If the words or ideas of this white paper are reproduced, extracted or otherwise used, please make a footnote that the source is from Moffett AI. In case of violation of the above statement, the company will investigate the relevant legal responsibilities.

Contents

How great things are made.....	03
The four great computing problems of our age.....	04
The processor brick wall.....	05
Sparsity in neural networks.....	06
Types of sparsity in neural networks.....	07
Why aren't neural networks sparse yet?.....	09
Sparsity becomes more promising.....	10
The unrealized potential of sparsity.....	11
Moffett AI and the path to dual sparsity.....	12
Co-designing hardware and software.....	14
A vision.....	15
Closing thoughts.....	16
About Moffett AI.....	17
The Sparse ecosystem.....	18
Patents.....	19
Published papers.....	20

How great things are made

Someone once asked Michelangelo how he carved the statue of David. “Was it difficult?”, they asked. “No,” he replied, “It was easy. I just chipped away the stone that didn’t look like David.”

When we think of building new things, we tend to think about adding materials together to create something larger and more complex. But in the case of David, the creative process was one of gradual subtraction, not addition.

There is a field of science built on applying this principle to how data is processed, called Sparse computing¹. The concept originated in the US in the 1950s² and is now led by large companies like NVIDIA and startups like Moffett AI. Sparse computing has become an important factor in modern chip design.

Sparsity is a smart way to process data. It is vital in applications where there are large amounts of data which need to be processed quickly, accurately and at affordable prices, consuming the least amount of energy.

Sparse computing makes it easier to process large quantities of data by removing the bits of data - and it turns out that there are always many of them - which don’t count. In summary, this vast field of Sparse computing is dedicated simply to how to remove data which is not useful so that the rest can be more easily processed.

The technique is becoming particularly important in the field of artificial intelligence. AI models, which are being used to model climate change, fold proteins, sequence DNA, drive Tesla vehicles, and translate speech and languages. They have grown so large that they are hard to run on normal computers. The largest ones, called “foundation models”, can cost millions of dollars in electricity just to develop, let alone operate. Some of them are so large that they have to be kept on specialized hardware. But by using sparsity to gradually remove unnecessary parameters, researchers are finding ways to make these models hundreds of times smaller – and thus usable at much more cost-effective manners for commercial applications.

¹ We capitalize the initial letter “S” in sparse to define the adjective. We use “Sparsity” as the noun, but “sparse” as the verb.

² The paper “Portfolio Selection” by Harry Markowitz of The Rand Corporation, was published in The Journal of Finance in March 1952

The four great computing problems of our age

We believe Sparse computing has an incredibly important role to play outside the laboratory, in the the real world.

Humanity today faces challenges from climate change and global health to wellbeing, safety, education and more. The United Nations Sustainable Development Goals³ provide a good summary of the main areas of focus and identify 17 major areas where work is required for humanity to address these challenges.

Computing plays a very big role in solving these challenges. At one end, modelling climate change and sequencing DNA gets better the more data there is, and so computers are processing larger and larger quantities of data in order to develop better insights. The more data they can process, the more accurate they get.

This data needs to be processed quickly. If you can process billions of images and data points in situation where time is critical, you can save lives.

At the other end, there are the vast server farms which house these computers, consuming more and more electricity to the extent that they are responsible for at least a few percentage points of the world's total energy consumption and a few percentage points of the world's total greenhouse gas emissions.

Data centers are a vital part of the solution, but also part of the problem: as they process more data, yet they so they consume more power.

³ The 2030 Agenda for Sustainable Development, adopted by all United Nations Member States in 2015, provides a shared blueprint for peace and prosperity for people and the planet, now and into the future. At its heart are the 17 Sustainable Development Goals (SDGs), which are an urgent call for action by all countries - developed and developing - in a global partnership. They recognize that ending poverty and other deprivations must go hand-in-hand with strategies that improve health and education, reduce inequality, and spur economic growth - all while tackling climate change and working to preserve our oceans and forests. <https://sdgs.un.org/goals>

That leaves four very clear problems to solve:

1. Increase speed - we need to be able to process more data than ever before in less time and decrease latency.
2. No loss of accuracy - we require computers which can process data with high precision.
3. Decrease cost - we can't simply add more and more processors and electricity, we need to make computing more affordable and reduce the total cost of ownership (TCO) from both capital and operation.
4. Decrease energy consumption - at current levels, our energy consumption is unsustainable and we need to do reduce it as much as possible.

Sparse computing helps all these four things. The more urgency we see in addressing the four compute problems, the more of a role there is for Sparse computing as a part of the solution.

The processor brick wall

The solution to most problems like this used to be to get smarter in how processor chips were designed and developed, specifically how many transistors could be engineered into a single chip.

In 1965, it became possible to create chips with 50 transistors, and it was predicted that this could grow to 1,000 by 1970. Intel made the prediction true with its first 2,300-transistor processor, but the number has since grown to 80 billion.

During this exponential growth, the number doubled approximately every 18 months, and the phenomenon became known as "Moore's Law". For over five decades, Moore's Law reliably provided more and more power but the laws of physics are now preventing further exponential growth and Moore's Law will cease to apply.

This "Moore's Wall" has been anticipated for some time. We saw that the CPU took processing capabilities only so far before the GPU (a graphics processing unit) was developed to handle the higher quantities of data which image processing required.

The GPU has triumphed since then. But it, in turn, is now running out of the capacity to keep up with demand.


Scientists and researchers are working with technologies such as Field Programmable Gate Arrays (FPGAs), Quantum Computing, neuromorphic

computing, memristors, graphene processors, nanotube transistors, soft machines and system-on-a-chip (SoC) devices to evolve processing performance. Sparse computing is one of the possible solutions. Moffett AI, began testing an FPGA Sparse processor in 2019 with public availability on Amazon AWS shortly after. Its Antoum compute platform is a Sparse processor, a system on a chip design which sees public availability in June 2022.


Sparsity in neural networks

Neural networks consist of a set of simulated neurons, analogous to the neurons which make up a biological brain, and a set of weighted connections between them. In these models, the behavior of a single neuron can be written as $f(x * w)$, where w is an ordered list, also called a vector, of N numbers which together represent the synaptic weights of a neuron. They describe the strength of the neuron's connections to N of its neighbors. Meanwhile, x is a vector of equal length which describes the activations of those neighbors.

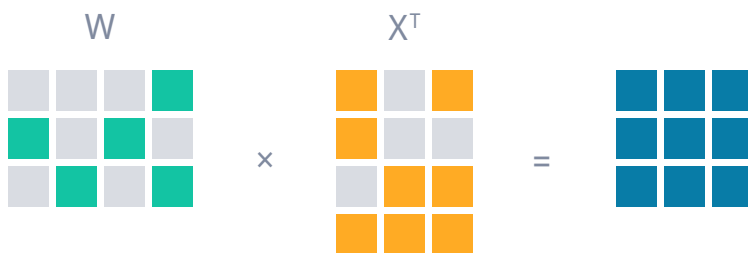
Simulating the behavior of this virtual neuron involves calculating a dot product between the two vectors and then applying the activation function f to the result.



A group of M of these neurons, all with the same number of synaptic connections, could be represented with an $M \times N$ matrix where each row is a separate neuron. In the same way, B different activations, each related to a different stimulus, could be stacked atop one another to form a $B \times N$ matrix. Using these groupings, we could compute the responses of a group of M neurons to B different inputs with a single matrix multiplication, $W * X^T$.



So far, we've assumed that every neuron has a connection to every element in the activation vector. But this is not generally the case, either in biological systems or artificial ones. It's much more common for a neuron to have a few strong synapses and let the rest be zero. Neuron activations tend to be sparse as well. This means that sparse matrices are a natural way to represent both the W and X matrices.



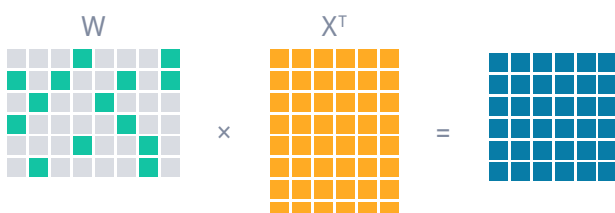
The diagram illustrates the multiplication of two sparse matrices, W and X^T , resulting in a dense matrix. Matrix W is a 4x4 grid with 6 green squares (non-zero elements) at positions (1,4), (2,1), (2,3), (3,2), (3,4), and (4,4). Matrix X^T is a 4x4 grid with 6 orange squares (non-zero elements) at positions (1,1), (1,4), (2,1), (3,2), (3,3), and (4,3). The resulting matrix is a 4x4 grid with 9 blue squares (all elements are non-zero).

Types of sparsity in neural networks

For early neural network researchers, the importance of sparsity was reinforced when they looked at the human brain. At the time, it was rare for people to stack more than a few layers of neurons atop one another and in these types of “shallow” networks, the total number of synapses would grow quadratically with the number of neurons. But this is not how synapses scale in the brain.

Synaptic sparsity


By some estimates, the human brain has 86 billion neurons and 150 trillion synapses⁴. These numbers imply that only 0.000005%⁵ of the possible connections between neurons are actually present. In other words, the connectivity of the brain is 99.999995% sparse. In this regime, the total number of synapses grows linearly with the number of neurons. Each biological neuron gets a fixed number of connections and this number doesn't change even as the total number of neurons increases. Researchers call this property synaptic sparsity.



The diagram illustrates the multiplication of a sparse matrix W and a dense matrix X^T . Matrix W is a 5x5 grid with 10 green squares (non-zero elements) at positions (1,2), (1,4), (2,1), (2,3), (3,2), (3,4), (4,1), (4,3), (5,2), and (5,4). Matrix X^T is a 5x5 grid with all orange squares (all elements are non-zero). The resulting matrix is a 5x5 grid with all blue squares (all elements are non-zero).


Activation sparsity

The human brain is not only sparse in synapses; it is also sparse in neuron activations. The energy consumed by a biological neuron is roughly proportional to the number of times it fires. So the fewer neurons that fire in the brain, the less energy it consumes. The brain uses this activation sparsity to save energy. By contrast, a simulated neuron as described above consumes the same amount of energy regardless of whether it fires or not. If its output is zero, that zero still gets multiplied with other numbers.



Dual sparsity

These two types of sparsity are complementary to one another. Activation sparsity allows signals to be routed through specific subsets of a network, while synaptic sparsity keeps those subsets small and efficient. Working together, they lead to much greater efficiency gains than would be possible if only one were being used. Researchers suspect that this “dual sparsity” is what permits the brain to be so efficient.



Neural network researchers of the 1990’s were aware of the benefits of sparsity and put a great deal of effort into sparsifying their models, with approaches such as weight magnitude regularization and magnitude pruning to achieve high levels of synaptic sparsity. These works show that sparsity was important even in the early days of AI.

4 Not all scientists will agree with these numbers but “The remarkable, yet not extraordinary, human brain as a scaled-up primate brain and its associated cost” authored by Suzana Herculano-Houzel and published in the Proceedings of the National Academy of Sciences of the United States of America, Vol 109 in June 2012, is a good reference.

5 We obtained this number as follows: $\text{sparsity} = \frac{\text{observed number of connections}}{\text{possible connections btwn } N \text{ neurons}} = \frac{150 \times 10^{12}}{N(N-1)/2} = \frac{150 \times 10^{12}}{(86 \times 10^9)(86 \times 10^9 - 1)/2} = 4 \times 10^{-8}$

Why aren't neural networks sparse yet?

However, sparse neural networks languished through the 2000's and early 2010's. The early 2000s were known as the AI Winter, a time when funding for neural network research dropped precipitously. Sparsity languished in the early 2010's, even though this was a time when there was a huge interest in and funding for AI. During this "AI Spring", as it came to be known, progress in other areas of AI occurred at a dizzying pace while sparsity stagnated.

It is possible that this was because there were so many other fruitful ways to improve models. First of all, there was better data. By the late 2000's, the internet had exploded in size which made it possible for researchers to construct massive datasets from publicly available data. Second, computing infrastructure grew much better. Not only did computers in general improve, but researchers found that they could massively accelerate their models by putting them on GPUs. A third important event was the rise of automatic differentiation (autodiff) frameworks like Theano, TensorFlow, and PyTorch. These frameworks made it easier to design new models, train them on specialized hardware, and run them on large datasets.


It's worth noting that neither the GPUs nor the autodiff frameworks were built with Sparse computing in mind. And so while they enabled big advances in model size and architecture, they made it very difficult for researchers to reap rewards from sparsity. As long as significant progress was happening in other areas, this was to remain the case. But as the 2010's drew to a close, questions of energy efficiency and the compute-vs-accuracy tradeoff became more pressing and sparsity became much more attractive.

Sparsity becomes more promising


By 2020, AI models and datasets were growing more quickly than the availability of compute power. We can take as an example three famous AI models: ResNet (2014) had 25 million parameters, BERT (2018) had 340 million parameters, and GPT-2 and 3 (2019 and 2020) had 1.5 and 175 billion parameters.

Meanwhile, the rate of improvement in chip technology was beginning to slow. Moore's Law, which had held steady for several decades, was beginning to break down as the sizes of transistors shrank to the limits set by physics. This led to a world in which computing power was increasingly scarce. In this world, the computational benefits of sparsity started to look very attractive⁶.

MOORE'S LAW IS NEAR ITS PHYSICAL LIMITS



CHIP CAPABILITY INCREASES SLOW DOWN



Steps toward sparsity

One of the early signs that a transition towards sparsity was underway was that academic publications referencing sparsity increased dramatically from 2018 onwards. Then, in 2020, NVIDIA released a chip called the A100 which featured a "sparsity processing unit" (SPU) with a 2x performance boost. In the same year, Google researchers took a first step towards adding sparsity support to Tensor Processing Units⁷ ("Sparse-TPU"). Since then, other companies, like Intel and Microsoft, have taken steps in the same direction⁸.


NUMBER OF PUBLICATIONS ON SPARSITY IN ACADEMIA



Sparsity is one of the most popular AI research fields

SPARSITY IN THE INDUSTRY

A milestone in AI 2.0: On May 15, 2020, Nvidia launched A100 with 2x Sparsity



The commercialization of sparsity has started, and it will continue to lead the future

⁶ "Things that deal with sparse parallelism," said Raja Koduri, Intel's head of chip architecture, "...will give rise to some new architectural ideas that are very different from what we are doing in vector-matrix, which is very mainstream right now." quoted in ZD Net <https://www.zdnet.com/article/intel-data-bandwidth-sparsity-are-the-two-biggest-challenges-for-ai-chips> in August 2020.

The unrealized potential of sparsity

While these recent developments are moving in the right direction, it is important to put them in context. Researchers have shown that many models can be pruned until they are well over 95% sparse without damaging performance. Naively, this would suggest that such models could be made twenty times smaller and more efficient by adding synaptic sparsity alone. And yet, existing speedups due to model sparsification are only of the order of 2x.

The hardware problem

So far, companies have just made incremental modifications to existing chip architectures and this has not been enough to unlock the order of magnitude gains that are available in theory. One reason that progress has been slow is that sparsity is a difficult hardware problem. Adding full sparsity support means representing matrices and vectors differently on hardware. It means structuring matrix multiplications differently. It means parallelizing computations in different ways. Many experts believe that AI chips need to be rebuilt from the ground up.

Startups

This task, which requires daring and flexibility, is well suited for startups. Indeed, some of the best work being done in this area is happening at small companies. Numenta, a Bay Area startup, recently demonstrated a custom chip with hardware support that runs a popular vision architecture 100 times faster than more traditional chips. Another company, NeuralMagic, offers model sparsification for shrinking foundation models so that they can run on laptop CPUs instead of expensive data center GPUs. But in order to realize the full potential of sparsity, the industry is going to need to design both hardware and software together. So far, only a few startups have tried to do this. One of the most interesting and ambitious of these companies is Moffett AI.

7 See “Sparse-TPU: adapting systolic arrays for sparse matrices” authored by Xin He and others in ICS’20: Proceedings of the 34th ACM International Conference on Supercomputing, June 2020 <https://dl.acm.org/doi/10.1145/3392717.3392751>



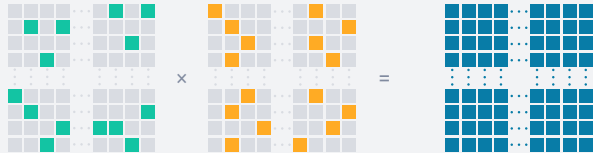
8 As just one example, in early 2022, Intel advertised an “Intel Neural Compressor” tool aimed at model sparsification <https://community.intel.com/t5/Blogs/Tech-Innovation/Artificial-Intelligence-AI/Quantizing-ONNX-Models-using-Intel-Neural-Compressor/post/1355237>

Moffett AI and the path to dual sparsity

Moffett AI was founded in 2018, and is focused specifically on building hardware and software infrastructure to support dual sparsity.

Dual sparsity

As discussed earlier, dual sparsity, refers to sparsity in both the weights and activations of neural networks. The diagram below gives an intuitive comparison of the differences between dense-dense operations, which the majority of AI chips use currently, dense-sparse operations, which some chips like the NVIDIA A100 offer, and “dual sparse” operations which Moffett AI supports. One thing to notice is that using dual sparsity permits researchers to evaluate the products of much larger matrices while using the same amount of memory, compute, and energy. In the upper image, Moffett’s approach leads to a substantial speedup.

NVIDIA V100 TENSOR CORE		Programmable Granularity: Small (for both AI & Graphics)
NVIDIA A100 TENSOR CORE		Bank-balanced Sparsity: Bank Size: 4 Support Sparsity up to 1/2
MOFFETT SPARSE TENSOR CORE		Programmable Granularity: Large (for AI) Bank-balanced dual Sparsity: Bank Size: 64 Support Sparsity up to 1/32

Four benefits of dual sparsity

The practical benefits of dual sparsity are fourfold: they include increased speed and decreased latency, no loss of accuracy, less energy consumption and lower cost thereby reducing TCO. To make these benefits more concrete, it is helpful to compare Moffett’s latest dual sparse chip, the Antoum, to its dense-dense and dense-sparse counterparts: the NVIDIA V100 and A100.



Antoum is an AI processor chip which combines hardware, software and algorithms to form an adaptive computing platform. Antoum may be built in to a range of devices to enable them for Sparse processing. It is available on a card which can be slotted into data center servers to easily and quickly make them Sparse-enabled.

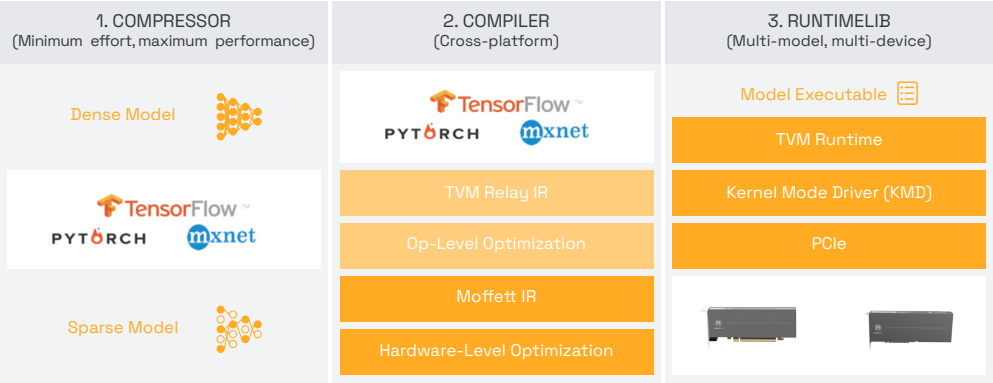
The first thing to notice is that it can run sparse models 32x and 8x faster than its two respective baselines. These gains are important because model latency matters a great deal in real-world settings. For example, a self-driving car needs to be able to process frames at a rate of at least 20 fps in order to avoid obstacles while moving at 60 miles per hour. As another example, internet users begin to lose interest in a webpage if its latency grows beyond a few dozen milliseconds, meaning that models used to filter information on these pages need to perform inference in about a millisecond.

The second benefit of dual sparse is that it allows us to run higher accuracy models for the same computational budget. Because inference is 8-32x faster on Antoum, it can run 8-32x larger models, which tend to be much more accurate.

A similar line of reasoning is behind the third benefit, which is less energy consumption. The energy that a chip consumes is roughly proportional to the number of mathematical operations it performs. Since dual sparse hardware ignores all the activations and weights that are set to zero, it saves the energy that would have been used to multiply them to get more zeros. These savings end up being between 8-32x as well. And since the cost of running a model scales with the amount of energy it uses, Antoum is 8-32x cheaper to operate.

Co-designing hardware and software

As the AI chip industry grows and evolves, Antoum and its baselines will soon be replaced by newer and better models. Given this observation, one of the things that really matters is potential for future improvements. Antoum is interesting because it is a first generation chip with a long roadmap of future improvements. In 2022, we have already achieved 32x sparsity which will be the best performance for the foreseeable future. Moffett’s research indicates that we could be achieving efficiency gains through dual sparsity which are as high as 100-1,000x in the not-too-distant future.



In order to realize these additional benefits, Moffett AI is focusing on co-designing sparse hardware and software together. Intuitively, this is a good idea because the best way to implement sparse operations in software depends on the chip being used. Meanwhile, the best dual sparse chip designs depend on the types of sparse models being run.

Take, for example, Moffett’s approach to running convolutional neural networks (CNNs) for computer vision. CNNs organize visual data along spatial and channel dimensions. Intuitively, spatial dimensions organize information according to where objects occur in a scene and channel dimensions do so according to what

objects occur in a scene. Given that objects, even common ones like wheels and eyes, occur infrequently throughout most images, Moffett's researchers realized that it is possible to make the channel dimension very sparse. Starting from this software observation, they adjusted the chip's physical design, allocating less processing power for the channel dimension.

This fertile cross-pollination between hardware and software engineering is rare at larger companies. It's much easier to achieve at small startups like Moffett, where the same researchers tend to be involved in both areas of development. While sometimes this occurs naturally, company structure also plays an important role. In the case of Moffett, for example, its founders chose each other with hardware and software co-design in mind. Two are hardware experts, one is a sparsity software expert, and the fourth has a background in both. Having joined together under a common banner, Moffett's founders have the unique combination of skills needed to make progress in this area.

A vision

Moore's Law in 1965 looked to the heavens to a world where the future would bring more and more transistors. Gordon Moore wrote: "integrated electronics will allow the advantages of electronics to be applied generally throughout society"⁹

and his vision, at that time, was about how economics and manufacturing technology, yields and physics could be mastered by chip makers.

"There remain," he concluded, "many significant problems for the electronics industry to solve in attempting to take advantage of this evolving technology to supply the rapidly increasing electronic requirements of the world."

The vision today in 2022, for the whole Sparse community, and particularly for Moffett AI, no longer lies with the technology, and certainly not with increases in yields.

The vision now is to reduce computation by the use of algorithms, which change the characteristics of the model itself, and combining algorithmic optimization with chip design, as Moffett AI is doing.

The spirit and ambition of Moore is alive and well, but the world of plenty in 1965 has now been replaced by a world of caution and prudence where it is now about doing less, and using sparse techniques to prune and thin out data. Over time, not only should all AI processing become sparse, but all processing could be sparse, and every computer could be a sparse computer.

⁹ From "Cramming more components onto integrated circuits" by Dr Gordon E Moore, Fairchild Semiconductor published in Electronics magazine, April 1965.

The mission at Moffett AI is to keep evolving the frontiers of AI processing using Sparse computing, and by extension to see that permeate to all processing, even beyond AI.

Which leads on to a new purpose, to power humanity's most demanding compute requirements with less energy and at lower cost, but faster and more accurately than ever before with Sparse computing.

We think Moore would have approved.¹⁰

Closing thoughts

Although the industry has started to put more time and energy into sparsity, there are many inefficiencies that have yet to be chiseled away. In coming years, we will need to adapt everything in AI, from chip design to low-level compilers and CUDA kernels to high-level autodiff frameworks, to better accommodate sparsity. Companies like Moffett AI are in a good position to lead this revolution. Perhaps the infrastructure they are building now will, in a few years, be running the most powerful AI models in the world.

¹⁰ Dr Gordon E Moore established the Gordon and Betty Moore Foundation in 2015 to make a "significant and positive impact in the world", tackling large, important issues at scale with the areas of interest including environmental conservation, scientific research, higher education, having observed changes in the natural world and from the dependency of all living species on the planet's health. <https://www.moore.org/about/founders-intent>

Appendix

About Moffett AI

Moffett AI (commonly referred to as Moffett) is the world leader in Sparse computing. Its Antoum Chip Architecture and S4, S10 and S30 Accelerators allow data centers to quickly and simply deploy sparse computing techniques in AI applications to increase processing speed, deliver higher accuracy, reduce processing cost and use less energy.

Moffett AI achieves these goals by combining hardware, software and algorithms on its computing platform using its 'Dual Sparse' proprietary technology to reduce the amount of data processed.

Moffett AI designs and produces the Sparse technology products which allow AI to solve some of humanity's greatest computing challenges. Our dream is to touch and improve every AI application and to embrace failure in the relentless pursuit of success.

Moffett AI was founded in 2018 and released its first product, the FPGA AI accelerator card in 2019. It owns over 30 patents related to Sparse computing.

Moffett AI comprises a team of computer scientists, chip designers, doers and dreamers united by and inspired by the philosophy of Sparse computing and a future where all data is sparsified. The team comes from backgrounds with Carnegie Mellon, Intel and Marvell.

The Sparse ecosystem

The Sparse ecosystem comprises fellow researchers, programmers, systems designers, data center operators, investors and others. It's time for the community to work together so that Sparse computing is acknowledged as the vital ingredient to achieve all four of the big goals (speed, accuracy, cost, and energy). Adding Sparse capabilities to existing computing facilities is now one-click easy.

AI cannot work without Sparse computing and Moffett is the world leader in this field. Moffett's Sparse computing platforms are needed to help meet humanity's most demanding and advanced computing requirements.

Patents

The patents below are assigned to Moffett AI

[System and method for knowledge-preserving neural network pruning](#)
11200497

[Method and system for hierarchical weight-sparse convolution processing](#)
11144823

[Method and system for balanced-weight sparse convolution processing](#)
11113601

[System and method for domain specific neural network pruning](#)
11068786

[Method and system for hierarchical weight-sparse convolution processing](#)
10970619

[Neural network acceleration and embedding compression systems and methods with activation sparsification](#)
10832139

Published papers

These are papers co-authored by Moffett founder, Ian E H Yen:


Title:	Sparse Progressive Distillation: Resolving Overfitting under Pretrain-and-Finetune Paradigm.
Authors:	Shaoyi Huang, Dongkuan Xu, Ian E.H. Yen, Sung-En Chang, Bingbing Li, Shiyang Chen, Mimi Xie, Hang Liu, Caiwen Ding. To Appear in Association for Computational Linguistics (ACL), 2022.
Title:	Rethinking Network Pruning -- under the Pre-train and Fine-tune Paradigm.
Authors:	Dongkuan Xu, Ian E.H. Yen, Jinxi Zhao, Zhibin Xiao. In Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL), 2021.
Title:	Minimizing FLOPs to Learn Efficient Sparse Representations.
Authors:	Authors: Biswajit Paria, Chih-Kuan Yeh, Ian E.H. Yen, Ning Xu, Pradeep Ravikumar, Barnabas Poczos. In International Conference on Learning Representations (ICLR), 2020.
Title:	Temporal Structure Mining for Weakly Supervised Action Detection.
Authors:	Tan Yu, Zhou Ren, Yuncheng Li, Ian E.H. Yen, Ning Xu, Junsong Yuan. In International Conference on Computer Vision (ICCV), 2019.
Title:	Efficient Global String Kernel with Random Features: Beyond Counting Substructures.
Authors:	Lingfei Wu, Ian En-Hsu Yen, Siyu Huo, Liang Zhao, Kun Xu, Liang Ma, Shouling Ji and Charu Aggarwal. In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2019.
Title:	Efficient Global String Kernel with Random Features: Beyond Counting Substructures.
Authors:	Lingfei Wu, Ian En-Hsu Yen, Siyu Huo, Liang Zhao, Kun Xu, Liang Ma, Shouling Ji and Charu Aggarwal. In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2019.
Title:	Scalable Global Alignment Graph Kernel Using Random Features: From Node Embedding to Graph Embedding.
Authors:	Lingfei Wu, Ian En-Hsu Yen, Zhen Zhang, Kun Xu, Liang Zhao, Xi Peng, Yinglong Xia and Charu Aggarwal. In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2019.

Title:	Representer Point Selection for Explaining Deep Neural Networks.
Authors:	Chih-Kuan Yeh*, Joon Sik Kim*, Ian E.H. Yen, and Pradeep Ravikumar. (equally contributed). In Advances in Neural Information Processing Systems (NerulPS), 2018.
Title:	MixLasso: Generalized Mixed Regression via Convex Atomic-Norm Regularization.
Authors:	Ian E.H. Yen, Wei-Cheng Lee, Kai Zhong, Sung-En Chang, Pradeep Ravikumar and Shou-De Lin. In Advances in Neural Information Processing Systems (NerulPS), 2018.
Title:	Word Mover's Embedding: From Word2Vec to Document Embedding.
Authors:	Lingfei Wu, Ian E.H. Yen, Kun Xu, Fangli Xu, Avinash Balakrishnan, Pin-Yu Chen, Pradeep Ravikumar and Michael J. Witbrock. In Empirical Methods in Natural Language Processing (EMNLP), 2018.
Title:	Loss Decomposition for Fast Learning in Large Output Spaces.
Authors:	Ian E.H. Yen, Satyen Kale, Felix Yu, Daniel Holtmann-Rice, Sanjiv Kumar, and Pradeep Ravikumar. In International Conference on Machine Learning (ICML), 2018.
Title:	Scalable Spectral Clustering Using Random Binning Features.
Authors:	Lingfei Wu, Pin-Yu Chen, Ian E.H. Yen, Fangli Xu, Yinglong Xia, and Charu Aggarwal. In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2018.
Title:	Random Warping Series: A Random Features Method for Time-Series Embedding.
Authors:	Lingfei Wu, Ian E.H. Yen, Jinfeng Yi, Fangli Xu, Qi Lei, and Michael J. Witbrock. In International Conference on Artificial Intelligence and Statistics (AISTATS), 2018.
Title:	PPDSparse: A Parallel Primal-Dual Sparse Method for Extreme Classification.
Authors:	Ian E.H. Yen, Xiangru Huang, Wei Dai, Pradeep Ravikumar, Inderjit S. Dhillon and Eric P. Xing. In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2017.
Title:	Latent Feature Lasso.
Authors:	Ian E.H. Yen, Wei-Cheng Li, Sung-En Chang, Arun S. Suggala, Shou-De Lin and Pradeep Ravikumar. In International Conference on Machine Learning (ICML), 2017.

Title:	Doubly Greedy Primal-Dual Coordinate Methods for Sparse Empirical Risk Minimization.
Authors:	Qi Lei, Ian E.H. Yen, Chao-Yuan Wu, Pradeep Ravikumar and Inderjit Dhillon. In International Conference on Machine Learning (ICML), 2017.
Title:	Greedy Direction Method of Multiplier for MAP Inference of Large Output Domain.
Authors:	Xiangru Huang, Ian E.H. Yen, Ruohan Zhang, Qixing Huang, Pradeep Ravikumar, and Inderjit Dhillon. In International Conference on Artificial Intelligence and Statistics (AISTATS), 2017.
Title:	Scalable Convex Multiple Sequence Alignment via Entropy-Regularized Dual Decomposition.
Authors:	Jiong Zhang, Ian E.H. Yen, Pradeep Ravikumar, and Inderjit Dhillon. In International Conference on Artificial Intelligence and Statistics (AISTATS), 2017.
Title:	Dual Decomposed Learning with Factorwise Oracles for Structural SVMs of Large Output Domain.
Authors:	Ian E.H. Yen, Xiangru Huang, Kai Zhong, Ruohan Zhang, Pradeep Ravikumar and Inderjit S. Dhillon. In Advances in Neural Information Processing Systems (NerulPS), 2016.
Title:	Revisiting Random Binning Features: Fast Convergence and Strong Parallelizability.
Authors:	Lingfei Wu*, Ian E.H. Yen*, Jie Chen and Rui Yan. (* equally contributed) In ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2016.
Title:	Large-scale Submodular Greedy Exemplar Selection with Structured Similarity Matrices.
Authors:	Dmitry Malioutov, Abhishek Kumar and Ian E.H. Yen. In Conference on Uncertainty in Artificial Intelligence (UAI), 2016.
Title:	PD-Sparse: A Primal and Dual Sparse Approach to Extreme Multiclass and Multilabel Classification.
Authors:	Ian E.H. Yen*, Xiangru Huang*, Kai Zhong, Pradeep Ravikumar and Inderjit S. Dhillon. (* equally contributed) In International Conference on Machine Learning (ICML), 2016.

Title:	A Convex Atomic-Norm Approach to Multiple Sequence Alignment and Motif Discovery.
Authors:	Ian E.H. Yen*, Xin Lin*, Jiong Zhang, Pradeep Ravikumar and Inderjit S. Dhillon. (* equally contributed) In International Conference on Machine Learning (ICML), 2016.
Title:	Scalable Exemplar Clustering and Facility Location via Augmented Block Coordinate Descent with Column Generation.
Authors:	Ian E.H. Yen, Dmitry Malioutov, and Abhishek Kumar. In International Conference on Artificial Intelligence and Statistics (AISTATS), 2016.
Title:	Sparse Linear Programming via Primal and Dual Augmented Coordinate Descent.
Authors:	Ian E.H. Yen, Kai Zhong, Cho-Jui Hsieh, Pradeep Ravikumar and Inderjit S. Dhillon. In Advances in Neural Information Processing Systems (NerulPS), 2015.
Title:	A Dual-Augmented Block Minimization Framework for Learning with Limited Memory.
Authors:	Ian E.H. Yen, Shan-Wei Lin, Shou-De Lin. In Advances in Neural Information Processing Systems (NerulPS), 2015.
Title:	A Convex Exemplar-based Approach to Dirichlet Process Mixture Models.
Authors:	Ian E.H. Yen, Jimmy Lin, Kai Zhong, Pradeep Ravikumar and Inderjit S. Dhillon. In International Conference on Machine Learning (ICML), 2015.
Title:	Tackling the Achilles Heel of Social Networks: Influence Propagation based Language Model Smoothing.
Authors:	Rui Yan, Ian E.H. Yen, Cheng-Te Li, Shiqi Zhao and Hu Xiaohua. In International World Wide Web Conference (WWW), 2015.
Title:	Constant Nullspace Strong Convexity and Fast Convergence of Proximal Methods under High-Dimensional Settings.
Authors:	Ian E.H. Yen, Cho-Jui Hsieh, Pradeep Ravikumar and Inderjit S. Dhillon. In Advances in Neural Information Processing Systems (NerulPS), 2014.
Title:	Sparse Random Feature as Coordinate Descent in Hilbert Space.
Authors:	Ian E.H. Yen, Ting-Wei Lin, Shou-De Lin, Pradeep Ravikumar and Inderjit S. Dhillon. In Advances in Neural Information Processing Systems (NerulPS), 2014.

-
- Title:** Proximal Quasi-Newton for Computationally Intensive L1-regularized M-estimators.
- Authors:** Kai Zhong, Ian E.H. Yen, Inderjit S. Dhillon and Pradeep Ravikumar. In Advances in Neural Information Processing Systems (NerulPS), 2014.
-
- Title:** Optimal Tests of Treatment Effects for the Overall Population and Two Subpopulations in Randomized Trials, using Sparse Linear Programming.
- Authors:** Michael Rosenblum, Han Liu, Ian E.H. Yen. In Journal of American Statistical Association (JASA) (Theory and Methodology), 2014.
-
- Title:** Indexed Block Coordinate Descent for Large-Scale Linear Classification with Limited Memory.
- Authors:** Ian E.H. Yen, Chun-Fu Chang, Ting-Wei Lin, Shan-Wei Lin, Shou-De Lin. In ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), 2013.
-
- Title:** On Convergence Rate of Concave-Convex Procedure.
- Authors:** Ian E.H. Yen, Nanyun Peng, Po-Wei Wang and Shou-de Lin. In NIPS Optimization for Machine Learning Workshop, 2012.
-
- Title:** Presenting Time-Evolving Activities Using Communication Archive Data.
- Authors:** Sheng-Jie Luo, Shih-Wen Huang, En-Hsu Yen, Li-Ting Huang, Bing-Yu Chen, Kwan-Liu Ma. In ACM SIGGRAPH ASIA (poster), 2011.
-
- Title:** Feature engineering and classifier ensemble for KDD Cup 2010.
- Authors:** H.F. Yu, H.Y. Lo, H.P. Hsieh, J.K. Lou, T.G. McKenzie, J.W. Chou, P.H. Chung, C.H. Ho, C.F. Chang, Y.H. Wei, J.Y. Weng, En-Hsu Yen, C.W. Chang, T.T. Kuo, Y.C. Lo, P.T. Chang, C. Po, C.Y. Wang, Y.H. Huang, C.W. Hung, Y.X. Ruan, Y.S. Lin, S.D. Lin, H.T. Lin, C.J. Lin. In Proceedings of the KDD Cup 2010 Workshop, July 2010. First-place winner report of KDD Cup 2010.
-



MOFFETT AI



<https://moffett.ai/>

Silicon Valley Office:
949 Sherwood Ave, Suite 200
Los Altos, CA 94022 USA
(650) 209-5845

Shenzhen Office:
24D Microsoft
Ketong Building,
Nashan District,
Shenzhen, China